

From Predictive Agreement to Geometric Agreement: Stability of Fisher Connections and Semantic Transport in Neural Representations

Guy Basson

`guy-shal.huji.ac.il@mail.huji.ac.il`

August 2026

Abstract

A neural hidden state is an ordinary vector, but through a probabilistic decoder it also parameterizes a predictive distribution. The predictively identifiable quotient therefore inherits a Fisher metric and the canonical family of Amari α -connections. This observation alone is classical and does not imply that models with similar cross-entropy possess similar local geometry. We make the missing implication precise. First, we prove a quantitative stability ladder: after a smooth cross-model alignment, C^1 closeness of predictive maps controls their Fisher metrics, C^2 closeness controls their α -connections and parallel transports, and C^3 closeness controls curvature. Second, we give a regular Bernoulli counterexample in which excess cross-entropy converges uniformly to zero while the pullback Fisher metric does not converge. Third, under a uniform Sobolev bound, we upgrade integrated KL convergence to C^2 convergence and obtain an explicit loss-to-transport rate. Finally, we introduce a semantic-sufficiency obstruction. The conditional Fisher variance of a semantic predictive effect is exactly the irreducible error of every vector field that depends only on the predictive state. This yields an orthogonal decomposition of cross-model transfer error into predictive insufficiency, field mismatch, and connection mismatch. The results clarify what would make geometric convergence across independently trained models mathematically natural, what pretraining alone cannot guarantee, and how parallel transport can be tested against ordinary vector reuse without presupposing that it wins.

1 Introduction

Modern neural networks compute with vectors while producing probability distributions. For an autoregressive language model, a hidden state h entering the output head determines

$$F(h) = p(\cdot \mid h) \in \Delta_N^\circ, \tag{1}$$

where Δ_N° is the interior of the categorical simplex over a vocabulary of size N . The hidden coordinate remains a point of $\mathbb{R}^d$. Nevertheless, its predictively visible component parameterizes a statistical model and inherits the Fisher geometry of the decoder.

This leads to a tempting but overly strong argument. Population cross-entropy is minimized by the true conditional distribution, so two sufficiently trained models should approach the same predictions. If both models represent the same statistical manifold, their Fisher geometries should agree. Semantic transformations should then be compared by parallel transport rather than by adding the same ambient vector at every context.

Each step requires qualification. Cross-entropy controls probability values on the data distribution, whereas Fisher metrics and affine connections depend on derivatives of the predictive map.

Hidden states also contain decoder-null directions, and a semantic transformation need not be determined by the next-token distribution. Finally, ordinary vector addition is not geometry-free: at a linear softmax head it is exactly the affine operation selected by the flat exponential connection. The scientific question is therefore not “vectors or manifolds,” but which connection, if any, makes a semantic effect transferable across contexts and models.

This paper formalizes that question. Our contributions are:

1. We state the exact predictive-naturality result that serves as the classical baseline: two identifiable parameterizations of the same predictive family preserve the Fisher metric, Amari-Chentsov tensor, every α -connection, parallel transport, curvature, and intrinsic semantic fields up to reparameterization.
2. We show that predictive loss convergence is strictly weaker than geometric convergence. A smooth, uniformly regular Bernoulli family converges uniformly in KL while its Fisher metric fails to converge at a fixed point.
3. We prove quantitative C^k stability bounds. On compact uniformly interior and uniformly identifiable families, C^2 predictive closeness controls the defect between aligned α -connections and the commutator of their parallel transports. C^3 closeness controls curvature.
4. We bridge population loss and geometry under an explicit regularity hypothesis. Integrated KL convergence plus a uniform H^s bound implies C^2 convergence by Sobolev interpolation, and hence convergence of α -parallel transport with a quantitative rate.
5. We prove a semantic conditional-variance decomposition. It separates failure of a semantic effect to descend to the predictive quotient from failure of a chosen connection and from genuine cross-model field mismatch.

The exact naturality statements and the analytic tools used in the proofs are classical. The differentiated contribution is their explicit quotient-aware synthesis into a predictive-map-to- α -transport stability theory and a semantic-sufficiency test for neural representations. We do not claim a first Fisher geometry of neural latent space, nor do we claim that parallel transport is universally better than vector addition.

2 Related work and novelty boundary

Information geometry equips regular statistical models with the Fisher metric, the Amari-Chentsov cubic tensor, and dual affine connections [2, 3, 6]. Invariance under sufficient statistics and the uniqueness of the Fisher and Amari-Chentsov tensors are classical [5, 8]. The distinction between affine flatness and Levi-Civita flatness is standard; global exponential or mixture coordinates do not force zero Riemannian curvature [7].

Pullback geometry in learned latent spaces is also established. Arvanitidis et al. [4] pull Fisher-Rao geometry through probabilistic decoders. Zhang [18] explicitly connects a language model’s hidden-to-predictive map, likelihood training, and pullback geometry. For softmax representations, Park et al. [12] develop an information-geometric account of probing and steering, while Wang and Zhao [15] pull the output Fisher metric through transformer layers and show large deviations from ambient Euclidean geometry. Thus neither probabilistic latent geometry nor Fisher-aware steering is new here.

Cross-model representation convergence has been studied from several directions. The Platonic Representation Hypothesis proposes convergence toward shared statistical structure [11]. Yu

et al. [17] align distinct neural latent geometries with pullback metrics and relative geodesic representations. Hu et al. [10] model contextual transformations as vector fields and report shared displacement structure across language-model families. Path dependence and holonomy have also been proposed as representation diagnostics [14]. These works motivate, but do not supply, the stability implication studied here.

Continuity of Fisher information is known to require control of both distributions and their derivatives; Rezakhani et al. [13] prove such bounds in classical and quantum settings. Quantitative relaxations of sufficient statistics based on bi-Lipschitz Fisher metrics are studied by Yamaguchi and Nozawa [16]. Our narrow target is different: after a comparison diffeomorphism, we propagate predictive-map closeness through every Amari α -connection to an explicit parallel-transport commutator bound, then derive the required derivative closeness from KL under Sobolev regularity.

3 Predictive statistical manifolds

3.1 The categorical simplex

Let

$$\Delta_N^\circ = \left\{ p \in \mathbb{R}^N : p_a > 0, \sum_{a=1}^N p_a = 1 \right\}. \quad (2)$$

For tangent vectors u, v, w satisfying $\sum_a u_a = \sum_a v_a = \sum_a w_a = 0$, define

$$g_p^\Delta(u, v) = \sum_{a=1}^N \frac{u_a v_a}{p_a}, \quad (3)$$

$$C_p^\Delta(u, v, w) = \sum_{a=1}^N \frac{u_a v_a w_a}{p_a^2}. \quad (4)$$

These are the Fisher metric and Amari-Chentsov tensor. A smooth predictive map

$$F : M \longrightarrow \Delta_N^\circ \quad (5)$$

induces

$$g_F = F^* g^\Delta, \quad C_F = F^* C^\Delta. \quad (6)$$

When F is an immersion, g_F is a Riemannian metric. In local coordinates $x = (x^1, \dots, x^m)$, writing $p_a = F_a(x)$,

$$(g_F)_{ij} = \sum_a \frac{\partial_i p_a \partial_j p_a}{p_a}, \quad (7)$$

$$(C_F)_{ijk} = \sum_a \frac{\partial_i p_a \partial_j p_a \partial_k p_a}{p_a^2}. \quad (8)$$

Raise the final index of C_F using g_F and define K_F by

$$g_F(K_F(X, Y), Z) = C_F(X, Y, Z). \quad (9)$$

We use the convention

$$\nabla_F^{(\alpha)} = \nabla_F^{\text{LC}} - \frac{\alpha}{2} K_F. \quad (10)$$

Thus $\alpha = 0$ is Levi-Civita, while $\alpha = 1$ and $\alpha = -1$ are the exponential and mixture connections on a regular exponential family.

3.2 Predictive quotients

Neural predictive maps are commonly nonidentifiable. Define

$$h \sim_F h' \iff F(h) = F(h'). \quad (11)$$

When the fibers are regular, the identifiable state space is the quotient $M_F = M/\sim_F$. All geometric statements below concern this quotient or a chosen horizontal subbundle on which the pullback metric is nondegenerate.

For a linear softmax decoder,

$$F(h) = \text{softmax}(Wh + b), \quad (12)$$

the null space is

$$\mathcal{N} = \{u : Wu \in \text{span}\{\mathbf{1}\}\}. \quad (13)$$

The quotient $\mathbb{R}^d/\mathcal{N}$ embeds smoothly into the simplex, and before quotienting

$$G(h) = W^\top (\text{diag } p_h - p_h p_h^\top) W, \quad \ker G(h) = \mathcal{N}. \quad (14)$$

This is the rigorous sense in which the predictively identifiable component of a hidden space is a statistical manifold. A finite set of naturally occurring hidden states remains a point cloud unless a smooth intervention chart or interpolation model is supplied.

3.3 Ordinary addition selects a connection

For the affine softmax family

$$p_h(a) = \exp(w_a^\top h + b_a - A(h)), \quad (15)$$

natural hidden coordinates are exponential-affine. Therefore constant coordinate vectors are $\nabla^{(1)}$ -parallel. More explicitly,

$$p_{h_r+h_q-h_p}(a) = \frac{p_{h_r}(a)p_{h_q}(a)/p_{h_p}(a)}{\sum_b p_{h_r}(b)p_{h_q}(b)/p_{h_p}(b)}. \quad (16)$$

Thus hidden vector arithmetic at the final affine head is exactly exponential transport, not the absence of geometry. Levi-Civita transport instead preserves Fisher lengths and angles, while mixture transport preserves the dual affine structure. Which connection best models a semantic operation is an empirical question.

4 Exact predictive naturality

We first state the exact result. It is classical reparameterization naturality and serves as a baseline rather than a novelty claim.

Theorem 4.1 (Exact predictive naturality). *Let $F_A : M_A \rightarrow \Delta_N^\circ$ and $F_B : M_B \rightarrow \Delta_N^\circ$ be smooth embeddings onto the same embedded predictive image P . Define*

$$\Phi = F_B^{-1} \circ F_A : M_A \longrightarrow M_B. \quad (17)$$

Then

$$\Phi^* g_B = g_A, \quad \Phi^* C_B = C_A. \quad (18)$$

For every $\alpha \in \mathbb{R}$,

$$d\Phi(\nabla_A^{(\alpha)} X Y) = \nabla_B^{(\alpha)} d\Phi X (d\Phi Y). \quad (19)$$

For every piecewise C^1 path γ ,

$$d\Phi_{\gamma(1)} P_\gamma^{A,\alpha} = P_{\Phi \circ \gamma}^{B,\alpha} d\Phi_{\gamma(0)}. \quad (20)$$

Curvature tensors intertwine, and holonomies of corresponding loops are conjugate.

Proof. Since $F_A = F_B \circ \Phi$, functoriality of pullback gives

$$\Phi^* g_B = \Phi^* F_B^* g^\Delta = (F_B \circ \Phi)^* g^\Delta = F_A^* g^\Delta = g_A, \quad (21)$$

and identically for C . A Riemannian isometry preserves the Levi-Civita connection. Preservation of g and C implies preservation of K , so (10) gives (19). Applying $d\Phi$ to an A -parallel field and using uniqueness of the linear transport ODE gives (20). Curvature naturality follows from its definition as the commutator of covariant derivatives, and loop holonomies are consequently conjugate. $\square$

Corollary 4.2 (Intrinsic semantic fields). *Let $S_t : P \rightarrow P$ be a smooth semantic flow with infinitesimal field*

$$Z_P(p) = \left. \frac{d}{dt} \right|_{t=0} S_t(p). \quad (22)$$

Define model fields X_i by $dF_i(X_i) = Z_P \circ F_i$. Then

$$d\Phi(X_A) = X_B \circ \Phi. \quad (23)$$

The condition is substantive: a context operation need not descend to P . If two hidden states have the same prediction before an intervention but distinct predictive effects afterward, then no intrinsic field Z_P exists on the predictive quotient. Section 8 quantifies this obstruction.

5 Prediction convergence does not imply geometric convergence

Let $P(\cdot | c)$ be the data conditional and $q(\cdot | c)$ a model. Population cross-entropy decomposes as

$$\mathcal{L}(q) - H_P(Y | C) = \mathbb{E}_C D_{\text{KL}}(P(\cdot | C) \| q(\cdot | C)). \quad (24)$$

The logarithmic score is strictly proper [9]. If models q_A, q_B have excess risks ϵ_A, ϵ_B , Pinsker's inequality and the triangle inequality yield

$$\mathbb{E}_C \|q_A(\cdot | C) - q_B(\cdot | C)\|_1 \leq \sqrt{2\epsilon_A} + \sqrt{2\epsilon_B}. \quad (25)$$

This makes output agreement natural under realizability and global population-risk optimization. It provides no derivative control.

Proposition 5.1 (Regular KL-to-geometry counterexample). *There exist smooth, uniformly interior, regular Bernoulli predictive families p_n, p_* on $[-1, 1]$ such that*

$$\sup_x D_{\text{KL}}(p_*(x) \| p_n(x)) = O(n^{-2}), \quad (26)$$

but their pullback Fisher metrics do not converge even pointwise at $x = 0$.

Proof. Let $p_*(x) = \text{Ber}(r(x))$ and $p_n(x) = \text{Ber}(r_n(x))$, where

$$r(x) = \frac{1}{3} + \frac{x}{20}, \quad r_n(x) = r(x) + \frac{1}{100n} \sin(nx). \quad (27)$$

All probabilities remain in a fixed compact subset of $(0, 1)$. Moreover,

$$r'_n(x) = \frac{1}{20} + \frac{1}{100} \cos(nx) \in [0.04, 0.06], \quad (28)$$

so every map is a regular immersion with a uniform rank bound. Uniform interiority and Taylor's theorem give $\sup_x D_{\text{KL}}(p_* \| p_n) = O(n^{-2})$. The Bernoulli pullback metric is

$$g(x) = \frac{r'(x)^2}{r(x)(1-r(x))} dx^2. \quad (29)$$

At $x = 0$, $r_n(0) = r(0) = 1/3$, but $r'_n(0) = 3/50$ and $r'(0) = 1/20$. Hence

$$g_n(0) = \frac{81}{5000} dx^2 \neq \frac{9}{800} dx^2 = g_*(0). \quad (30)$$

□

This example rules out an unconditional theorem of the form “cross-entropy convergence implies Fisher convergence.” The oscillations have bounded first derivative but unbounded higher regularity, anticipating the Sobolev condition in Section 7.

There is a second nonidentifiability. For any hidden diffeomorphism ψ ,

$$H' = \psi \circ H, \quad F' = F \circ \psi^{-1} \quad (31)$$

produces exactly the same predictions while changing ambient Euclidean coordinates. Pullback Fisher geometry handles this gauge covariantly, but loss alone cannot identify a privileged hidden Euclidean metric.

6 Stability of Fisher geometry and transport

We now quantify how derivative-level predictive agreement propagates to geometry. We state the result on a common Euclidean chart. For two latent manifolds related by a comparison diffeomorphism Φ , replace the second map by $\tilde{p} = F_B \circ \Phi$.

Let $K \Subset U \subset \mathbb{R}^m$ and let $p, \tilde{p} : U \rightarrow \Delta_N^\circ$. For an integer k , use the componentwise norm

$$\|p\|_{C^k(K)} = \max_{a, |\beta| \leq k} \sup_{x \in K} |\partial^\beta p_a(x)|. \quad (32)$$

Assumption 6.1 (Uniform regularity). For constants $\eta, \lambda, B_2 > 0$,

$$p_a(x), \tilde{p}_a(x) \geq \eta, \quad (33)$$

$$\|p\|_{C^2(K)} + \|\tilde{p}\|_{C^2(K)} \leq B_2, \quad (34)$$

$$\lambda_{\min}(g_p(x)), \lambda_{\min}(g_{\tilde{p}}(x)) \geq \lambda \quad (35)$$

for every a and $x \in K$.

The probability floor excludes the singular simplex boundary. The spectral floor excludes decoder-null directions after quotienting and prevents inverse-metric instability.

Theorem 6.2 (C^2 stability of α -transport). *Under Assumption 6.1, fix $A > 0$. There is a constant*

$$C_\Gamma = C_\Gamma(m, N, A, \eta^{-1}, \lambda^{-1}, B_2) \quad (36)$$

such that, uniformly for $|\alpha| \leq A$,

$$\begin{aligned} & \|g_p - g_{\tilde{p}}\|_{C^1(K)} + \|C_p - C_{\tilde{p}}\|_{C^0(K)} \\ & + \|\Gamma_p^{(\alpha)} - \Gamma_{\tilde{p}}^{(\alpha)}\|_{C^0(K)} \leq C_\Gamma \|p - \tilde{p}\|_{C^2(K)}. \end{aligned} \quad (37)$$

Let $\gamma : [0, 1] \rightarrow K$ be piecewise C^1 , with Euclidean length L_γ . If

$$M_\Gamma = \max\{\|\Gamma_p^{(\alpha)}\|_{C^0}, \|\Gamma_{\tilde{p}}^{(\alpha)}\|_{C^0}\}, \quad (38)$$

where the norm is the induced bilinear operator norm, then

$$\|P_\gamma^{p,\alpha} - \tilde{P}_\gamma^{p,\alpha}\|_{\text{op}} \leq L_\gamma e^{M_\Gamma L_\gamma} C_\Gamma \|p - \tilde{p}\|_{C^2(K)}. \quad (39)$$

Proof. Equations (7) and (8) are finite sums of rational functions of p and its first derivatives. Their denominators are controlled by η . Differentiating (7) once uses at most second derivatives, giving the first two terms of (37) by the mean-value theorem on the compact regularity class.

For completeness, the lower-index α -coefficients have the direct probability-coordinate expression

$$\Gamma_{ij,k}^{(\alpha)} = \sum_a \left[\frac{\partial_{ij} p_a \partial_k p_a}{p_a} - \frac{1 + \alpha}{2} \frac{\partial_i p_a \partial_j p_a \partial_k p_a}{p_a^2} \right]. \quad (40)$$

Thus their difference is $O(\|p - \tilde{p}\|_{C^2})$. Raising the final index requires inverse metrics. The identity

$$g_p^{-1} - g_{\tilde{p}}^{-1} = g_p^{-1} (g_{\tilde{p}} - g_p) g_{\tilde{p}}^{-1} \quad (41)$$

and the spectral floor λ give inverse-metric stability, proving (37).

Along γ , transport solves the linear ODE

$$\dot{V}(t) = -\Gamma^{(\alpha)}(\gamma(t))[\dot{\gamma}(t)]V(t). \quad (42)$$

Duhamel's formula bounds the difference of the two fundamental solutions by the integral of the coefficient defect multiplied by their growth factors. Gronwall's inequality yields (39). $\square$

Corollary 6.3 (Cross-model transport commutator). *Let $F_A : M_A \rightarrow \Delta_N^\circ$, $F_B : M_B \rightarrow \Delta_N^\circ$, and let $\Phi : M_A \rightarrow M_B$ be a C^2 local diffeomorphism. Suppose F_A and $F_B \circ \Phi$ satisfy Assumption 6.1 and*

$$\|F_A - F_B \circ \Phi\|_{C^2(K)} \leq \varepsilon. \quad (43)$$

Then

$$\begin{aligned} & \|\text{d}\Phi_{\gamma(1)} P_\gamma^{A,\alpha} - P_{\Phi \circ \gamma}^{B,\alpha} \text{d}\Phi_{\gamma(0)}\|_{\text{op}} \\ & \leq L_\gamma e^{M_\Gamma L_\gamma} C_\Gamma \varepsilon, \end{aligned} \quad (44)$$

after using the aligned tangent norms. At $\varepsilon = 0$, this reduces to Theorem 4.1.

Proof. The connection induced by $F_B \circ \Phi$ is the pullback of the connection induced by F_B . Its transport is conjugated by $\text{d}\Phi$. Apply Theorem 6.2 on M_A . $\square$

Theorem 6.4 (C^3 curvature stability). *Assume the hypotheses of Theorem 6.2, replace the C^2 bound by*

$$\|p\|_{C^3(K)} + \|\tilde{p}\|_{C^3(K)} \leq B_3, \quad (45)$$

and retain the probability and spectral floors. Then

$$\|R_p^{(\alpha)} - R_{\tilde{p}}^{(\alpha)}\|_{C^0(K)} \leq C_R \|p - \tilde{p}\|_{C^3(K)}, \quad (46)$$

where C_R depends on $m, N, A, \eta^{-1}, \lambda^{-1}, B_3$. Levi-Civita sectional curvatures also converge on corresponding two-planes whose Gram determinants are uniformly bounded below.

Proof. Differentiating the argument for (37) gives C^1 control of the Christoffel symbols. Insert it into

$$R_{ijk}^\ell = \partial_i \Gamma_{jk}^\ell - \partial_j \Gamma_{ik}^\ell + \Gamma_{ir}^\ell \Gamma_{jk}^r - \Gamma_{jr}^\ell \Gamma_{ik}^r. \quad (47)$$

Sectional curvature is a quotient. The same estimates control its metric numerator, while the uniform Gram bound controls its denominator. $\square$

Remark 6.5 (Why the assumptions matter). The constants diverge near $p_a = 0$ because the Fisher tensors contain p_a^{-1} and p_a^{-2} . They also diverge near rank loss because raising indices requires g^{-1} . These are geometric instabilities, not proof artifacts. In large-vocabulary models, a useful extension would replace a uniform tokenwise probability floor by moment or coarse-graining conditions with vocabulary-independent constants.

7 From cross-entropy to transport under Sobolev regularity

Theorem 6.2 assumes derivative-level agreement. We now give sufficient conditions under which population KL convergence implies it.

Let Ω be a compact m -dimensional smooth manifold or a bounded Sobolev-extension domain. Let μ have density ρ with respect to volume, satisfying $\rho \geq \rho_0 > 0$. Define

$$\mathcal{K}(p, \tilde{p}) = \int_{\Omega} D_{\text{KL}}(p(x) \| \tilde{p}(x)) \, d\mu(x). \quad (48)$$

Theorem 7.1 (Risk-to-transport convergence). *Let $p, \tilde{p} : \Omega \rightarrow \Delta_N^\circ$ satisfy the probability and spectral floors of Assumption 6.1. Suppose*

$$p, \tilde{p} \in H^s(\Omega; \mathbb{R}^N), \quad s > 2 + \frac{m}{2}, \quad (49)$$

and

$$\|p\|_{H^s} + \|\tilde{p}\|_{H^s} \leq B. \quad (50)$$

For every r with

$$2 + \frac{m}{2} < r < s, \quad (51)$$

there is a constant C such that, for every bounded-length path γ ,

$$\|P_\gamma^{p, \alpha} - P_\gamma^{\tilde{p}, \alpha}\|_{\text{op}} \leq C B^{r/s} \mathcal{K}(p, \tilde{p})^{\frac{s-r}{2s}}. \quad (52)$$

Equivalently, the transport error is $O(\mathcal{K}^\beta)$ for every

$$0 < \beta < \frac{s - 2 - m/2}{2s}. \quad (53)$$

Proof. Pinsker's inequality and $\rho \geq \rho_0$ give

$$\begin{aligned} \|p - \tilde{p}\|_{L^2(\text{dvol})}^2 &\leq \rho_0^{-1} \int_{\Omega} \|p - \tilde{p}\|_1^2 \, \text{d}\mu \\ &\leq 2\rho_0^{-1} \mathcal{K}(p, \tilde{p}). \end{aligned} \quad (54)$$

Sobolev interpolation [1] yields

$$\|p - \tilde{p}\|_{H^r} \leq C \|p - \tilde{p}\|_{L^2}^{1-r/s} \|p - \tilde{p}\|_{H^s}^{r/s}, \quad (55)$$

and therefore

$$\|p - \tilde{p}\|_{H^r} \leq C B^{r/s} \mathcal{K}(p, \tilde{p})^{(s-r)/(2s)}. \quad (56)$$

Since $r > 2 + m/2$, Sobolev embedding gives $H^r \hookrightarrow C^2$. Apply Theorem 6.2. $\square$

Corollary 7.2 (Two-model geometric convergence). *Suppose two model sequences $p_{A,n}, p_{B,n}$ on a shared intervention chart converge in population cross-entropy to the same conditional map p_* , and all three maps satisfy uniform versions of the hypotheses of Theorem 7.1. Then, for every admissible β ,*

$$\|P_{\gamma}^{A,n,\alpha} - P_{\gamma}^{B,n,\alpha}\|_{\text{op}} \leq C(\epsilon_{A,n}^{\beta} + \epsilon_{B,n}^{\beta}), \quad (57)$$

where $\epsilon_{i,n} = \mathbb{E}D_{\text{KL}}(p_* \| p_{i,n})$. With a cross-model alignment Φ_n , the same statement holds for the transport commutator in (44) after replacing $p_{B,n}$ by $p_{B,n} \circ \Phi_n$.

If $s > 3 + m/2$, the same argument combined with Theorem 6.4 gives curvature convergence for every

$$0 < \beta < \frac{s - 3 - m/2}{2s}. \quad (58)$$

The uniform Sobolev hypothesis is strong but transparent. It rules out the high-frequency behavior in Proposition 5.1. For language models it should be tested on low-dimensional continuous intervention charts, such as soft-prompt paths or controlled activation subspaces, rather than assumed on the discrete set of tokenized contexts.

8 The semantic-sufficiency obstruction

Even perfect agreement of predictive geometry does not guarantee that a semantic operation is a field on that geometry. We now separate this representational obstruction from connection mismatch.

Let $F : \mathcal{H} \rightarrow P$ be a regular predictive quotient, let H be a random hidden state with law μ , and let $X_T(H) \in T_H \mathcal{H}$ describe a semantic operation T . Its observable infinitesimal predictive effect is

$$Z_T(H) = \text{d}F_H X_T(H) \in T_{F(H)} P. \quad (59)$$

Assume square integrability in the Fisher metric. Since P is embedded in the simplex, tangent-valued conditional expectations may be taken fiberwise in the ambient zero-sum space; conditioning on $F(H) = p$ fixes both $T_p P$ and g_p .

Theorem 8.1 (Semantic conditional-variance obstruction). *For a measurable, square-integrable field $Y : P \rightarrow TP$, define*

$$\mathcal{E}_T(Y) = \mathbb{E} \|Z_T(H) - Y(F(H))\|_{g_{F(H)}}^2. \quad (60)$$

Let

$$\bar{Z}_T(p) = \mathbb{E}[Z_T(H) \mid F(H) = p]. \quad (61)$$

Then

$$\begin{aligned} \mathcal{E}_T(Y) &= \mathbb{E}\|Z_T - \bar{Z}_T(F(H))\|_g^2 \\ &\quad + \mathbb{E}\|\bar{Z}_T(F(H)) - Y(F(H))\|_g^2. \end{aligned} \quad (62)$$

Consequently,

$$\inf_Y \mathcal{E}_T(Y) = \mathbb{E} \text{Var}_g(Z_T \mid F). \quad (63)$$

The semantic effect descends almost surely to a predictive field if and only if this conditional variance is zero. Smooth descent additionally requires a smooth version of $\bar{Z}_T$.

Proof. Square-integrable tangent-valued random variables form a Hilbert direct integral with inner product

$$\langle U, V \rangle = \mathbb{E}[g_{F(H)}(U, V)]. \quad (64)$$

The fields measurable with respect to $\sigma(F)$ form a closed subspace. Conditional expectation is the orthogonal projection onto that subspace. Equivalently, decompose

$$\begin{aligned} Z_T - Y(F) &= (Z_T - \bar{Z}_T(F)) \\ &\quad + (\bar{Z}_T(F) - Y(F)). \end{aligned} \quad (65)$$

The two terms are orthogonal after conditioning on F , which proves (62) and (63). $\square$

Corollary 8.2 (Connection-restricted decomposition). *Let $\mathcal{C}_\alpha$ be any closed class of predictive fields selected by an α -connection, for example fields satisfying $\nabla^{(\alpha)}Y = 0$ on a specified simply connected chart or fields with a bounded covariant-derivative residual. Then*

$$\begin{aligned} \inf_{Y \in \mathcal{C}_\alpha} \mathcal{E}_T(Y) &= \underbrace{\mathbb{E} \text{Var}_g(Z_T \mid F)}_{\text{predictive insufficiency}} \\ &\quad + \underbrace{\inf_{Y \in \mathcal{C}_\alpha} \mathbb{E}\|\bar{Z}_T - Y\|_g^2}_{\text{connection mismatch}}. \end{aligned} \quad (66)$$

Corollary 8.3 (Cross-model semantic transfer). *Let P_A, P_B be predictive quotients related by a Fisher isometry $\Phi : P_A \rightarrow P_B$ on paired samples. Write $Q_i = F_i(H_i)$ for the random predictive state of model i . Define*

$$Y_{A \rightarrow B}(p_B) = d\Phi_{\Phi^{-1}(p_B)} \bar{Z}_A(\Phi^{-1}(p_B)). \quad (67)$$

Then

$$\begin{aligned} \mathbb{E}\|Z_B - Y_{A \rightarrow B}(Q_B)\|_{g_B}^2 &= \underbrace{\mathbb{E}\|Z_B - \bar{Z}_B(Q_B)\|_{g_B}^2}_{\text{model-}B \text{ quotient obstruction}} \\ &\quad + \underbrace{\mathbb{E}\|\bar{Z}_B - d\Phi \bar{Z}_A\|_{g_B}^2}_{\text{cross-model field mismatch}}. \end{aligned} \quad (68)$$

No choice of α can reduce the first term of (66). Therefore a large conditional variance falsifies the claim that the next-token predictive manifold is a sufficient semantic state for T before exponential, Levi-Civita, mixture, or learned- α transport is compared.

9 Operational tests and training implications

The theorems suggest a hierarchy of increasingly strong cross-model agreement tests:

$$\begin{aligned}
\text{Level 0} & : \text{output KL agreement,} \\
\text{Level 1} & : \text{Fisher metric agreement,} \\
\text{Level 2} & : \alpha\text{-connection and transport agreement,} \\
\text{Level 3} & : \text{semantic-field equivariance,} \\
\text{Level 4} & : \text{curvature and holonomy agreement.}
\end{aligned} \tag{69}$$

Agreement at one level does not imply the next without the regularity and sufficiency conditions identified above.

9.1 Transport commutation score

For an estimated alignment Φ and sampled paths and tangent vectors, define

$$\mathcal{D}_{\text{PT}}^{(\alpha)}(A, B; \Phi) = \mathbb{E}_{\gamma, v} \frac{\|\text{d}\Phi_{\gamma(1)} P_{\gamma}^{A, \alpha} v - P_{\Phi\gamma}^{B, \alpha} \text{d}\Phi_{\gamma(0)} v\|}{\|v\|}. \tag{70}$$

Theorem 4.1 makes this score zero under exact predictive equivalence, while Corollary 6.3 bounds it under approximate equivalence. A practical study should test whether (70) predicts held-out steering transfer or task generalization beyond output KL and representation-only similarities such as CKA.

9.2 Selecting the connection

At a linear softmax head, $\alpha = 1$ is ordinary vector reuse, $\alpha = 0$ is Fisher-metric-compatible Levi-Civita transport, and $\alpha = -1$ is mixture transport. A semantic field X_T can be evaluated through

$$\mathcal{R}_T(\alpha) = \mathbb{E} \|\nabla^{(\alpha)} X_T\|_g^2. \tag{71}$$

The correct interpretation is comparative:

- exponential transport winning supports context-independent natural-coordinate vectors for the tested operation;
- Levi-Civita winning supports preservation of predictive Fisher geometry;
- mixture or an intermediate α winning supports probabilistic affine structure without privileging metric compatibility;
- every connection failing, together with a large conditional variance in (63), falsifies sufficiency of the predictive quotient;
- every connection failing with small conditional variance motivates a richer, operation-specific connection or nonparallel field model.

9.3 Training regimes

The risk-to-transport theorem indicates why derivative regularity matters. A causal training study can compare cross-entropy-only models with matched-loss models trained using

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{CE}} + \lambda \widehat{\text{Var}}_g(dFX_T | F) + \mu \|\nabla^{(\alpha)} \bar{Z}_T\|_g^2 + \nu \mathcal{R}_{\text{smooth}}. \quad (72)$$

The first regularizer encourages the operation to descend to the predictive quotient. The second selects connection-relative consistency. The final term controls the derivative oscillations excluded by Theorem 7.1. Behavioral improvement, rather than a smaller geometric regularizer alone, must be the primary outcome.

A no-GPU pilot can use frozen small models and one- or two-dimensional soft-prompt charts. Finite differences estimate F , ∂F , and $\partial^2 F$; the model’s output head provides exact vocabulary probabilities. The main comparisons are output KL, metric defect, the transport score (70), semantic conditional variance, and held-out intervention transfer.

10 Limitations

Continuous charts. Tokenized contexts are discrete. Derivatives and connections require an explicit continuous chart, such as soft prompts, activation interventions, or a justified local manifold model.

Boundary and rank. Softmax probabilities are positive but can be extremely small. Uniform tokenwise floors are mathematically clean and practically strong. The Fisher metric can also be severely ill-conditioned. Quotienting, regularization, coarse-graining, or horizontal-subspace analysis is necessary near rank loss.

Different tokenizers. The exact theorem assumes a common outcome space. A tokenizer map generally acts as a Markov morphism. Fisher geometry is preserved only under a statistically sufficient or reversible identification; otherwise it contracts [5].

Population versus empirical loss. Equation (24) concerns population risk, realizability, and global optimization. Finite empirical loss introduces estimation and optimization errors not treated here.

Alignment is not identified automatically. The stability theorem evaluates a proposed Φ . Constructing a statistically identifiable alignment from finite observations, especially across different latent dimensions, is a separate problem. One may instead work on a shared low-dimensional intervention chart.

No universal superiority claim. Parallel transport compares tangent vectors at different points once a connection is selected. It does not make ambient vector addition invalid. At the final affine softmax head, addition is exactly exponential-affine transport. The preferred connection can depend on the operation, layer, and task.

11 Conclusion

Idealized pretraining makes convergence of conditional output distributions natural because excess cross-entropy is expected KL divergence. It does not by itself make convergence of Fisher metrics, affine connections, or curvature natural. Geometry depends on derivatives, and a regular oscillatory counterexample separates these notions.

Under explicit smoothness and identifiability assumptions, the stronger conclusion becomes provable. Predictive maps that converge in C^2 have convergent Amari α -connections and parallel transports; C^3 convergence controls curvature. Integrated KL plus a uniform Sobolev bound supplies the missing derivative convergence and yields a quantitative loss-to-transport theorem. The semantic conditional-variance obstruction then determines whether the predictive quotient is rich enough to support the operation at all.

The resulting research hypothesis is precise and falsifiable: independently trained models that estimate the same smooth predictive system should, after quotienting and alignment, exhibit small α -transport commutators and equivariant semantic fields. Failure can be localized to inadequate regularity, predictive insufficiency, field mismatch, or connection mismatch. This is a stronger and more useful claim than saying that latent spaces are “really” probabilistic rather than linear.

Code and manuscript availability

The manuscript source, model-free predictive-geometry implementation, CPU tests, and experimental protocols are available at

<https://github.com/guybass/probalistic-latent-spaces>.

References

- [1] Robert A. Adams and John J. F. Fournier. *Sobolev Spaces*, volume 140 of *Pure and Applied Mathematics*. Academic Press, Amsterdam, 2 edition, 2003. ISBN 978-0-12-044143-3.
- [2] Shun-ichi Amari. *Information Geometry and Its Applications*, volume 194 of *Applied Mathematical Sciences*. Springer, Tokyo, 2016. doi: 10.1007/978-4-431-55978-8.
- [3] Shun-ichi Amari and Hiroshi Nagaoka. *Methods of Information Geometry*, volume 191 of *Translations of Mathematical Monographs*. American Mathematical Society and Oxford University Press, Providence, RI, 2000. ISBN 978-0-8218-0531-2.
- [4] Georgios Arvanitidis, Miguel González-Duque, Alison Pouplin, Dimitrios Kalatzis, and Soren Hauberg. Pulling back information geometry. In *Proceedings of the 25th International Conference on Artificial Intelligence and Statistics*, volume 151 of *Proceedings of Machine Learning Research*, pages 4872–4894. PMLR, 2022. URL <https://proceedings.mlr.press/v151/arvanitidis22b.html>.
- [5] Nihat Ay, Jürgen Jost, Hông Vân Lê, and Lorenz Schwachhöfer. Information geometry and sufficient statistics. *Probability Theory and Related Fields*, 162:327–364, 2015. doi: 10.1007/s00440-014-0574-8.
- [6] Nihat Ay, Jürgen Jost, Hông Vân Lê, and Lorenz Schwachhöfer. Parametrized measure models. *Bernoulli*, 24(3):1692–1725, 2018. doi: 10.3150/16-BEJ910.

- [7] Ovidiu Calin and Constantin Udriste. *Geometric Modeling in Probability and Statistics*. Springer Monographs in Mathematics. Springer, Cham, 2014. doi: 10.1007/978-3-319-07779-6.
- [8] N. N. Cencov. *Statistical Decision Rules and Optimal Inference*, volume 53 of *Translations of Mathematical Monographs*. American Mathematical Society, Providence, RI, 1982. doi: 10.1090/mmono/053.
- [9] Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007. doi: 10.1198/016214506000001437.
- [10] Zhimin Hu, Lanhao Niu, and Sashank Varma. Language models represent and transform concepts with shared geometry. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*. PMLR, 2026. URL <https://arxiv.org/abs/2607.04525>.
- [11] Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. Position: The platonic representation hypothesis. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 20617–20642. PMLR, 2024. URL <https://proceedings.mlr.press/v235/huh24a.html>.
- [12] Kiho Park, Todd Nief, Yo Joong Choe, and Victor Veitch. The information geometry of softmax: Probing and steering. In *Proceedings of the 43rd International Conference on Machine Learning*, 2026. URL <https://arxiv.org/abs/2602.15293>.
- [13] A. T. Rezakhani, M. Hassani, and S. Alipour. Continuity of the quantum fisher information. *Physical Review A*, 100(3):032317, 2019. doi: 10.1103/PhysRevA.100.032317.
- [14] Vasileios Sevetlidis and George Pavlidis. Gauge-invariant representation holonomy. In *The Fourteenth International Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2601.21653>.
- [15] Sihan Wang and Jiayi Zhao. FishBack: Pullback fisher geometry for optimal activation steering in transformers, 2026.
- [16] Kaori Yamaguchi and Hiraku Nozawa. On statistics which are almost sufficient from the viewpoint of the fisher metrics. *Information Geometry*, 7(2):543–553, 2024. doi: 10.1007/s41884-024-00160-1.
- [17] Hanlin Yu, Berfin Inal, Georgios Arvanitidis, Søren Hauberg, Francesco Locatello, and Marco Fumero. Connecting neural models latent geometries with relative geodesic representations. In *Advances in Neural Information Processing Systems*, volume 38, 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/a3ed50e08a0d016ee5ee1148a8a392c4-Abstract-Conference.html.
- [18] Yifan Zhang. A markov categorical framework for language modeling, 2025. Version 5 revised May 2026.